

This notebook finds the features from `listings.csv` that are most correlated with the rating. (See end of notebook for list of these features.)

```
In [1]: # TODO: better preprocessing of binary features
```

```
In [2]: import pandas as pd
import numpy as np
import sklearn
import sklearn.preprocessing
import re
import pickle
```

```
In [3]: train = pd.read_csv('data/listings.csv')

train.shape
```

```
/Users/ChentianJiang/miniconda3/envs/cs317/lib/python3.7/site-packages/IPython/c
ore/interactiveshell.py:2785: DtypeWarning: Columns (43,87,88) have mixed types.
Specify dtype option on import or set low_memory=False.
  interactivity=interactivity, compiler=compiler, result=result)
```

```
Out[3]: (50914, 96)
```

In [4]: `train.head()`

Out[4]:

	id	listing_url	scrape_id	last_scraped	name	summary
0	2515	<a href="https://www.airbnb.com/rooms/2515">https://www.airbnb.com/rooms/2515</a>	20180806171147	2018-08-07	Stay at Chez Chic budget room #1	Step into our artistic spacious apartment and ...
1	2539	<a href="https://www.airbnb.com/rooms/2539">https://www.airbnb.com/rooms/2539</a>	20180806171147	2018-08-07	Clean & quiet apt home by the park	Renovated apt home in elevated building
2	2595	<a href="https://www.airbnb.com/rooms/2595">https://www.airbnb.com/rooms/2595</a>	20180806171147	2018-08-07	Skylit Midtown Castle	Find your romantic getaway to this beautiful ...
3	3330	<a href="https://www.airbnb.com/rooms/3330">https://www.airbnb.com/rooms/3330</a>	20180806171147	2018-08-07	++ Brooklyn Penthouse Guestroom ++	This is a spacious clean, furnished master bedroom...
4	3647	<a href="https://www.airbnb.com/rooms/3647">https://www.airbnb.com/rooms/3647</a>	20180806171147	2018-08-07	THE VILLAGE OF HARLEM....NEW YORK !	NaN

5 rows × 96 columns

```
In [5]: # global vars
```

```
# drop useless columns
DROP_COLS = set(["scrape_id",
                 "last_scraped",
                 "thumbnail_url",
                 "medium_url",
                 "picture_url",
                 "xl_picture_url",
                 "host_thumbnail_url",
                 "host_picture_url",
                 "host_total_listings_count",
                 "host_has_profile_pic",
                 "calendar_last_scraped",
                 "availability_30",
                 "availability_60",
                 "availability_90",
                 "availability_365",
                 "first_review",
                 "last_review"])

LABEL = 'review_scores_rating' # TODO: is this the overall rating?
```

```
In [6]: train_label = train[LABEL]
train.drop(LABEL, axis=1, inplace=True)
train.shape
```

```
Out[6]: (50914, 95)
```

```
In [7]: train.drop(DROP_COLS, axis=1, inplace=True)
train.shape
```

```
Out[7]: (50914, 78)
```

```
In [8]: # separate out the different *types* of columns
train_type_dict = train.columns.to_series().groupby(train.dtypes).groups
train_type_dict.keys()
```

```
Out[8]: dict_keys([dtype('int64'), dtype('float64'), dtype('O')])
```

```
In [9]: def get_cat_cols(data_df, cat_colnames):
        # find categorical/character columns with <= 10 unique values
        # (features with too many categories is hard to work with)
        cols_cat = []

        for col in cat_colnames:
            if len(data_df[col].unique()) <= 10: # 3 to include binary values as well as NA/NaN/Null
                cols_cat.append(col)
        return cols_cat
```

```
In [10]: def create_data_dict(data_df, type_dict, label_df):
        int_data = data_df[type_dict[np.dtype('int')]]
        float_data = data_df[type_dict[np.dtype('float')]]
        num_data = pd.concat([int_data, float_data], axis=1)

        cat_colnames = get_cat_cols(data_df, type_dict[np.dtype('object')])
        cat_data = data_df[cat_colnames]
        return {'num': num_data, 'cat': cat_data, 'label': label_df}
```

```
In [11]: train_dict = create_data_dict(train, train_type_dict, train_label)

        train_dict['num'].shape, train_dict['cat'].shape, train_dict['label'].shape
```

```
Out[11]: ((50914, 23), (50914, 20), (50914,))
```

```
In [12]: DROP_COLS_NUM = set()
        DROP_COLS_CAT = set()
```

```
In [13]: # drop numerical columns with a high ratio (50%) of missing values

        train_num_missing_cols = train_dict['num'].columns[train_dict['num'].isnull().any()
        ()]
        for col in train_num_missing_cols:
            if (train_dict['num'][col].isnull().sum() >= 0.5*train.shape[0]):
                DROP_COLS_NUM.add(col)

        DROP_COLS_NUM
```

```
Out[13]: {'host_acceptance_rate', 'square_feet'}
```

```
In [14]: # doesn't make sense to predict overall review with other reviews --> drop review
        -related numerical columns

        train_num_review_cols = set([col for col in train_dict['num'].columns if re.searc
        h('review', col)])
        DROP_COLS_NUM = DROP_COLS_NUM.union(train_num_review_cols)

        DROP_COLS_NUM
```

```
Out[14]: {'host_acceptance_rate',
        'number_of_reviews',
        'review_scores_accuracy',
        'review_scores_checkin',
        'review_scores_cleanliness',
        'review_scores_communication',
        'review_scores_location',
        'review_scores_value',
        'reviews_per_month',
        'square_feet'}
```

```
In [15]: train_dict['num'].drop(DROP_COLS_NUM, axis=1, inplace=True)
        train_dict['num'].shape
```

```
Out[15]: (50914, 13)
```

```
In [16]: # fill in the numerical missing values with the column median
        # TODO: better ways to impute missing numerical values?
        def fill_na(num_df):
            missing_cols = num_df.columns[num_df.isnull().any()]
            for col in missing_cols:
                num_df[col].fillna(num_df[col].median(), inplace=True)

        fill_na(train_dict['num'])
```

```
In [17]: # drop categorical columns with a high ratio (50%) of missing values

train_cat_missing_cols = train_dict['cat'].columns[train_dict['cat'].isnull().any()
()]
for col in train_cat_missing_cols:
    if(train_dict['cat'][col].isnull().sum() >= 0.5*train.shape[0]):
        DROP_COLS_CAT.add(col)

DROP_COLS_CAT
```

```
Out[17]: {'jurisdiction_names', 'license'}
```

```
In [18]: train_dict['cat'].drop(DROP_COLS_CAT, axis=1, inplace=True)
train_dict['cat'].shape

/Users/ChentianJiang/miniconda3/envs/cs317/lib/python3.7/site-packages/pandas/co
re/frame.py:3697: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stabl
e/indexing.html#indexing-view-versus-copy
errors=errors)
```

```
Out[18]: (50914, 18)
```

```
In [19]: train_dict['cat'] = train_dict['cat'].astype('str') # consistent type
```

```
In [20]: # one-hot encode categorical features

train_dict['cat'] = pd.get_dummies(train_dict['cat']) # this also encodes NaN/Nul
l/NA etc. as its own category
train_dict['cat'].shape
```

```
Out[20]: (50914, 57)
```

```
In [21]: # normalize (put into range [0,1]) numerical features
train_num = sklearn.preprocessing.normalize(train_dict['num'], axis=0)
train_num = pd.DataFrame(train_num, columns=train_dict['num'].columns)

# standardize (center to the mean and scale to unit variance) numerical features
train_num = sklearn.preprocessing.scale(train_num, axis=0)
train_num = pd.DataFrame(train_num, columns=train_dict['num'].columns)
```

```
In [22]: # put together numerical and categorical data (separated from labels)
train_cat = train_dict['cat'].astype('int')
train_data = pd.concat([train_num, train_cat], axis=1)

train_label.shape, train_data.shape
```

```
Out[22]: ((50914,), (50914, 70))
```

```
In [23]: # impute missing label value with median
# TODO: better impute method?
train_dict['label'].fillna(train_dict['label'].median(), inplace=True)
```

```
In [24]: # choose features that are relatively more correlated with the label

corrDict = dict()
for col in train_data.columns:
    train_data[col].dtype
    corr = abs(train_data[col].corr(train_label)) # "computes pairwise correlati
on of columns, excluding NA/null values"
    if corr > 0.02: # most features are barely correlated with the label
        corrDict[col] = corr

# save chosen features
with open('corrDict.pickle', 'wb') as f:
    pickle.dump(corrDict, f)
len(corrDict) # number of features with correlation above a certain threshold
```

Out[24]: 19

```
In [25]: corrDict # top correlated (with label) features with correlation above some thres
hold
```

```
Out[25]: {'id': 0.04697486373512992,
'host_id': 0.03647428871160532,
'accommodates': 0.032057733695012615,
'calculated_host_listings_count': 0.026567959389538032,
'longitude': 0.023159218757749976,
'beds': 0.03612508486889648,
'host_is_superhost_f': 0.14965465554289106,
'host_is_superhost_t': 0.14973725471109212,
'host_identity_verified_f': 0.02716568689308185,
'host_identity_verified_t': 0.027181575202442035,
'neighbourhood_group_cleansed_Brooklyn': 0.03238031945441595,
'neighbourhood_group_cleansed_Manhattan': 0.027002073378334608,
'room_type_Entire home/apt': 0.03422772761577601,
'room_type_Private room': 0.03287318388260937,
'instant_bookable_f': 0.060519618291476764,
'instant_bookable_t': 0.060519618291476764,
'cancellation_policy_flexible': 0.021684854503989923,
'cancellation_policy_moderate': 0.03486031433140535,
'cancellation_policy_strict_14_with_grace_period': 0.04845450263351716}
```